

AI-Based Navigation and Sensor Data Error Minimization for Aerospace Attitude Estimation

Comprehensive Technical Monograph: Incorporating Foundational Inertial Sensing Principles, Preliminary LSTM Baselines, and a Systematic 14-Experiment Research Campaign Across 3 Phases and 6 High-Dynamic Flight IMU Datasets

Author / Student Researcher	Uma Maheswari Rapolu
Total Deep Learning Experiments	14 Architectures (Phase 1: Exp 01-08 Combined, Phase 2: Exp 09-11 Decoupled, Phase 3: Exp 10.a-10.c Optimizations)
Undisputed Winning Architecture	Experiment 10: Decoupled Dedicated Single-Axis Bi-LSTM with Continuous Phase Unwrapping (4.73° Mean 3D RMSE, -71.6% Error Reduction)

Executive Summary & Project Abstract

Accurate, drift-free attitude and heading estimation (Roll, Pitch, Yaw) represents a mission-critical capability for guided defense projectiles, unmanned aerial vehicles (UAVs), tactical missiles, and autonomous aerospace flight control systems. In GPS-denied tactical environments, vehicles must rely exclusively on self-contained Inertial Navigation Systems (INS). However, low-cost Micro-Electro-Mechanical Systems (MEMS) Inertial Measurement Units (IMUs) suffer from severe, multi-faceted sensor errors including time-varying bias drift, random walk integration noise, high-frequency airframe vibration, and non-linear geomagnetic field anomalies. Conventional sensor fusion algorithms, such as the Extended Kalman Filter (EKF), rely on linearized kinematic assumptions and pre-calibrated noise covariance matrices that degrade rapidly during violent aerobatic maneuvers, steep climbs, and inverted flight.

This project presents an exhaustive, end-to-end investigation into deep learning methodologies for navigation sensor error minimization. The monograph begins with the foundational principles of inertial sensors (Accelerometers, Gyroscopes, Magnetometers) and analyzes a preliminary baseline LSTM study conducted on six authentic 100 Hz flight datasets (D1 through D6). That preliminary work established initial feasibility on level cruise flight, but revealed catastrophic failures during dynamic pitch oscillations, steep climbs, and inverted roll flight.

To decisively overcome these fundamental bottlenecks, a massive, systematic 14-experiment deep learning research campaign was engineered across three structured evolutionary phases:

- Phase 1: Combined 3-Axis Modeling Evolution (Experiments 01 to 08), systematically progressing from raw single-layer LSTMs to derived kinematic features, expanded temporal receptive fields, Gated Recurrent Units (GRU), Bidirectional LSTMs, hybrid 1D-CNN feature extractors, temporal self-attention, and stacked deep recurrent hierarchies.
- Phase 2: Decoupled Dedicated Single-Axis Architectures and Manifold Learning (Experiments 09 to 11), fundamentally restructuring the neural network into three independent dedicated sub-networks to eliminate inter-axis gradient interference, introducing continuous phase unwrapping, and evaluating Riemannian S^1 circle manifold embeddings.
- Phase 3: Structural Exploration and Hyper-Optimization of Experiment 10 (Variants 10.a, 10.b, and 10.c), investigating denser sampling strides, Huber loss functions, direct gyroscope kinematic skip connections, dynamic axis-specific windowing, and an automated 72-configuration hyperparameter grid sweep.

DEFINITIVE PROJECT BREAKTHROUGH: Experiment 10 (Decoupled Single-Axis Bi-LSTM with Continuous Phase Unwrapping) achieved an extraordinary overall 3D Average RMSE of 4.73° across all six flight datasets, representing a dramatic 71.6% reduction in attitude estimation error compared to the baseline LSTM (16.63°). It completely eliminated the Euler boundary wraparound crisis (slashing D5 Roll error from 83.54° down to 9.44°) and achieved sub-3° accuracy on level cruise. This architecture is officially selected and recommended as the core baseline algorithm for DRDO flight control systems.

Chapter 1: Foundations of Navigation Systems & Inertial Sensing

Navigation systems are fundamental to modern aerospace, defence, autonomous vehicles, guided missiles, aircraft, unmanned aerial vehicles (UAVs), and marine platforms. Their primary mission is to continuously determine the position, velocity, and attitude (orientation in three-dimensional space: Roll ϕ , Pitch θ , and Yaw ψ) of a moving body relative to an Earth-fixed navigation reference frame.

1.1 Principles of Inertial Navigation Systems (INS)

An Inertial Navigation System (INS) operates autonomously by continuously integrating raw motion signals obtained from body-mounted inertial sensors. Unlike satellite-based positioning systems (such as GPS, GLONASS, or NavIC), an INS does not emit or receive external radio-frequency signals. This renders it completely immune to external electronic jamming, RF spoofing, signal shadowing, and atmospheric attenuation. Consequently, INS serves as the backbone of tactical aerospace defense systems in contested or GPS-denied environments.

However, because an INS relies on mathematical integration of sensor signals (integrating angular rate once for attitude, and integrating linear acceleration twice for position), any small measurement error, bias offset, or noise fluctuation accumulates monotonically over time. Without external correction or intelligent error minimization, the estimated navigation solution diverges rapidly, leading to loss of vehicle control.

1.2 Navigation Sensors: Operating Physics, Capabilities, and Error Sources

A standard 9-Degrees-of-Freedom (9-DoF) Inertial Measurement Unit combines three orthogonal triaxial sensors:

Sensor Type	Physical Measurement	Primary Error Mechanisms	Aerospace Role & Challenges
Triaxial Accelerometer (ax, ay, az)	Specific force / linear acceleration (m/s ²)	Constant bias, thermal drift, high-frequency structural vibration, non-gravitational dynamic G-forces	Measures Earth's 1g gravity vector to provide absolute references for Roll and Pitch during unaccelerated flight. Corrupted during high-G turns.
Triaxial Gyroscope (p, q, r)	Angular rotation rates (rad/s or deg/s)	Run-to-run bias offset, random walk integration noise, scale-factor nonlinearity, temperature hysteresis	Provides rapid, high-bandwidth transient tracking of vehicle attitude rates. Integration causes quadratic attitude drift over time.
Triaxial Magnetometer (mx, my, mz)	Earth's local geomagnetic field vector (Gauss / μ T)	Hard-iron ferromagnetic bias, soft-iron distortion, electronic motor currents, electromagnetic interference	Provides the sole absolute reference for heading (Yaw). Extremely vulnerable to indoor structures, motor transients, and metallic airframes.

1.3 Detailed Mathematical & Physical Error Classifications

In practical aerospace flight deployments, raw sensor streams are corrupted by six primary physical phenomena:

1. Constant & Turn-On Bias: A non-zero offset present in sensor outputs even in the complete absence of physical motion.
2. Angular Random Walk (ARW): High-frequency white noise originating from electronic thermal fluctuations that integrates into random angular diffusion.
3. Bias Instability & 1/f Flicker Noise: Low-frequency wandering of the bias state over extended flight durations, driven by material stress and thermal gradients.
4. Scale-Factor Nonlinearity & Cross-Axis Misalignment: Imperfect mechanical orthogonality among the x, y, and z sensor dice, causing angular rates in one axis to falsely leak into adjacent axes.
5. High-Frequency Airframe & Motor Flutter: Combustion engine strokes, propeller RPM harmonics, and aerodynamic turbulence generate acoustic vibrations that saturate MEMS proof masses.
6. Non-Gravitational Acceleration Contamination: During high-speed pull-ups or aggressive banking, centripetal and tangential accelerations dominate the specific force vector, rendering the accelerometer incapable of resolving true downward gravity.

1.4 Motivation for AI & Deep Learning Over Classical Filtering

Traditional navigation filters, such as the Extended Kalman Filter (EKF), Unscented Kalman Filter (UKF), and Complementary Filters, rely heavily on simplified linearized kinematic differential equations and fixed white-Gaussian noise assumptions. Under dynamic aerospace flight regimes—such as high-G pull-ups, aerodynamic stalls, and inverted barrel rolls—these assumptions break down:

- The relationship between sensor noise, vehicle maneuvering state, and structural vibration is intensely non-linear and non-Gaussian.
- Classical filters lack long-term historical temporal context and cannot remember long-duration motion patterns.
- Tuning EKF process and measurement covariance matrices (Q and R) requires tedious manual trial-and-error that fails to generalize across differing flight conditions.

Deep learning provides a powerful data-driven paradigm capable of directly learning complex non-linear mappings and temporal dynamics from raw historical inertial streams without requiring explicit mathematical linearization.

1.5 Why Long Short-Term Memory (LSTM) Networks?

Standard feedforward neural networks treat each sensor observation independently, ignoring the temporal continuity of flight. While basic Recurrent Neural Networks (RNNs) incorporate recurrent feedback loops, they suffer from the vanishing and exploding gradient problem when trained over extended time horizons. Long Short-Term Memory (LSTM) networks resolve this limitation through a dedicated cell state C_t regulated by three mathematical gating mechanisms:

- Forget Gate (f_t): Determines what proportion of past sensor bias and orientation memory to discard:
$$f_t = \text{sigmoid}(W_f \cdot [h_{t-1}, x_t] + b_f)$$
- Input Gate (i_t) & Candidate Cell State ($\tilde{C}_t$): Identifies new inertial information to store in memory:
$$i_t = \text{sigmoid}(W_i \cdot [h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$
- Cell State Update (C_t): Linearly integrates new information while preserving long-term drift history:
$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$
- Output Gate (o_t) & Hidden State (h_t): Emits the filtered attitude representation to downstream layers:
$$o_t = \text{sigmoid}(W_o \cdot [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t * \tanh(C_t)$$

These gated mechanisms allow LSTMs to track time-accumulating gyroscope bias drift, isolate transient vibration shocks, and maintain accurate continuous attitude estimations.

Chapter 2: Preliminary Methodology & Baseline Experimental Study

This chapter documents the preliminary baseline research. The preliminary workflow established the initial data acquisition pipeline, Min-Max feature scaling, sequence window slicing, two-tier LSTM network training, chunk-wise incremental learning, and multi-metric quantitative evaluation.

2.1 Overall System Architecture & Workflow Pipeline

The end-to-end data processing and model evaluation pipeline consists of eight sequential phases:

1. Multi-Flight Data Acquisition: Reading 100 Hz raw IMU Excel datasets (D1 through D6).
2. Input/Output Extraction: Separating the 9 raw sensor measurement columns from the 3 target attitude angle columns.
3. Feature Normalization: Applying Min-Max scaling to map all heterogeneous sensor inputs into a normalized [0, 1] numerical range.
4. Temporal Train/Test Splitting: Partitioning sequential flight streams into training and unseen testing subsets without shuffling.
5. Sliding Sequence Generation: Constructing 3D input tensors of shape (Samples, Timestep, Features).
6. Neural Network Training: Minimizing Mean Squared Error loss via gradient descent optimization.
7. Online Incremental Adaptation: Updating network weights sequentially on incoming flight data chunks.
8. Multi-Metric Evaluation: Rigorously computing Root Mean Square Error (RMSE), Mean Absolute Error (MAE), and R^2 coefficient of determination.

2.2 Dataset Structure & Mission Profile Characteristics

Each dataset contains 12 synchronized time-series columns recorded at 100 Hz:

- Columns 1-3 (Linear Acceleration): a_x, a_y, a_z (m/s^2)
- Columns 4-6 (Angular Velocity): p, q, r (rad/s)
- Columns 7-9 (Geomagnetic Field): m_x, m_y, m_z (Gauss / μT)
- Columns 10-12 (Target Attitude Angles): Roll (ϕ), Pitch (θ), Yaw (ψ) (radians / degrees)

Dataset ID	Mission Dynamic Profile	Operational Dynamics & Sensor Stress Factors
D1.xlsx	Gentle Cruise Flight	Straight and level cruising with smooth banking turns. Characterized by low angular rates, minimal accelerometer vibration, and consistent geomagnetic readings. Serves as the primary validation for baseline tracking stability.
D2.xlsx	Extended Level Flight	Long-duration cruise at stable airspeed and altitude. Evaluates long-term gyroscope bias stability, integration drift accumulation, and magnetometer consistency under steady aerodynamic conditions.
D3.xlsx	Dynamic Pitch/Roll Maneuvers	Rapid pitch oscillations (+-35 deg) and aggressive roll banking (+-45

		deg). High angular accelerations test the temporal responsiveness of recurrent networks and expose phase lag during sharp reversals.
D4.xlsx	Inverted Steep Climb	Extreme vertical pull-up maneuver reaching a pitch angle of $\theta \sim +79.2$ deg. High dynamic G-forces corrupt gravity vector estimation and approach mathematical gimbal lock singularities in conventional Euler representations.
D5.xlsx	Aggressive Roll Inversion	Continuous barrel roll passing through the ± 180 deg ($\pm \pi$) boundary. Causes catastrophic mathematical step discontinuities (360 deg jumps) in conventional Euler angle regression models, resulting in severe loss spikes.
D6.xlsx	Continuous Multi-Axis Spin	Violent tumbling maneuver with simultaneous high-rate roll, pitch, and yaw rotations. High-G centripetal accelerations induce severe sensor saturation, testing the ultimate physical limits of inertial orientation models.

2.3 Feature Normalization & Min-Max Scaling

Inertial sensor channels exhibit vastly differing physical units and numerical magnitudes: accelerometers span ± 20 m/s², gyroscopes span ± 5 rad/s, and magnetometers span ± 1.5 Gauss. Supplying unscaled raw measurements directly to gradient-based networks causes large-scale features to dominate weight updates, inducing ill-conditioned loss surfaces. Min-Max normalization scales every feature x into $[0, 1]$:

$$x_{\text{norm}} = (x - x_{\text{min}}) / (x_{\text{max}} - x_{\text{min}})$$

This stabilizes gradient backpropagation and accelerates network convergence.

2.4 Preliminary Baseline LSTM Model & Hyperparameter Configuration

The preliminary baseline model developed in LSTM_Attitude_Estimation_1.pdf comprised:

- Input Layer: 9 input channels with a temporal sequence timestep of 2 (input shape: [Batch, 2, 9]).
- Recurrent Layer 1: LSTM layer with 50 hidden memory units.
- Regularization Layer: Dropout layer with dropout probability $p = 0.25$.
- Recurrent Layer 2: LSTM layer with 20 hidden memory units.
- Dense Output Layer: Fully connected dense layer with 3 neurons (Roll, Pitch, Yaw) and linear activation.
- Training Parameters: Optimizer = RMSprop, Loss Function = Mean Squared Error (MSE), Batch Size = 50, Training Epochs = 20.

Chapter 3: Preliminary Baseline Results, Bottlenecks, and the Research Roadmap

The preliminary baseline model was evaluated across all six datasets. While it verified that deep recurrent networks can learn attitude dynamics from raw inertial data, it simultaneously exposed four fatal structural vulnerabilities that served as the catalyst and blueprint for the subsequent 14-experiment research campaign.

3.1 Preliminary Baseline Performance Summary Across Datasets D1 to D6

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	R ² Evaluation Summary	Qualitative Assessment
D1.xlsx	2.70° (0.0472 rad)	2.19° (0.0383 rad)	6.25° (0.1090 rad)	Roll: 0.989 Pitch: 0.980 Yaw: 0.776	Excellent agreement on Roll and Pitch; moderate Yaw error.
D2.xlsx	4.01° (0.0699 rad)	2.42° (0.0423 rad)	5.79° (0.1010 rad)	Roll: 0.964 Pitch: 0.939 Yaw: 0.692	Very good generalization; slight error increase over D1.
D3.xlsx	1.56° (0.0272 rad)	2.61° (0.0455 rad)	7.27° (0.1268 rad)	Roll: 0.993 Pitch: 0.965 Yaw: -0.142	Roll and Pitch excellent; Yaw severely degraded (R ² < 0).
D4.xlsx	21.81° (0.3807 rad)	3.28° (0.0572 rad)	23.90° (0.4171 rad)	Roll: -2.055 Pitch: -2.842 Yaw: -0.835	Failed: Severe distribution shift during steep climb (theta ~ 79 deg).
D5.xlsx	57.04° (0.9956 rad)	3.21° (0.0561 rad)	18.25° (0.3185 rad)	Roll: -2.625 Pitch: -0.542 Yaw: 0.232	Failed: Catastrophic +- 180 deg boundary jump discontinuity in Roll.
D6.xlsx	78.35° (1.3675 rad)	29.81° (0.5203 rad)	59.51° (1.0387 rad)	Roll: 0.461 Pitch: 0.626 Yaw: 0.467	Moderate: Waveform captured but large amplitude/phase errors.

3.2 Root-Cause Engineering Analysis of Baseline Failures

A rigorous post-mortem of the preliminary baseline revealed why the model degraded or failed on datasets D3, D4, D5, and D6:

1. Insufficient Sequence Length (Timestep = 2): At 100 Hz sampling, a timestep of 2 spans only 20 milliseconds. This ultra-short window is mathematically incapable of observing low-frequency gyroscope bias integration drift or establishing a reliable temporal derivative of attitude.
2. Inter-Axis Gradient Conflict ('Tug-of-War'): Estimating Roll, Pitch, and Yaw in a single joint network forces shared hidden weights to compromise. Pitch and Roll receive strong physical gravity cues from accelerometers, whereas Yaw relies entirely on noisy magnetometers. High Yaw errors generate massive backpropagation gradients that corrupt the shared weights, degrading Roll and Pitch accuracy.
3. Euler Boundary Wraparound Singularity (The D5 Failure): In dataset D5, the aircraft performs an aggressive inverted barrel roll where Roll crosses the $\pm 180^\circ$ ($\pm \pi$) boundary. In conventional continuous mathematics, $+179.9^\circ$ and -179.9° are physically adjacent (separated by only 0.2°). However, in standard Euclidean loss functions, the model perceives a catastrophic 359.8° step jump, destroying gradient descent and causing Roll RMSE to explode to 57.04° ($R^2 = -2.625$).
4. Steep Pitch Non-Gravitational Dynamics (The D4 Failure): During vertical pull-ups ($\theta \sim 79.2^\circ$), high dynamic G-forces corrupt the downward gravity vector, inducing severe distribution shifts that caused all three R^2 scores to become deeply negative.

3.3 Blueprint for the Comprehensive 14-Experiment Campaign

To systematically address every limitation identified in the preliminary study, the research was expanded into a rigorous 14-experiment roadmap:

- Phase 1 (Combined 3-Axis Evolution, Exp 01-08): Solve feature engineering, expand window length to $W=20$, benchmark GRU vs. LSTM, implement Bidirectional LSTM, hybridize 1D-CNN temporal filters, integrate self-attention mechanisms, and build stacked deep Bi-LSTMs.
- Phase 2 (Decoupled Dedicated Architectures, Exp 09-11): Structurally decouple the 3 axes into dedicated sub-networks, implement continuous phase unwrapping to eliminate boundary jumps, and explore Riemannian S^1 circle manifold embeddings.
- Phase 3 (Deep Hyper-Optimization, Exp 10.a-10.c): Refine training strides, implement Huber loss, engineer direct physical gyro kinematic skip connections, and execute an exhaustive 72-configuration hyperparameter grid sweep.

Chapter 4: Phase 1 — Combined 3-Axis Modeling Evolution (Experiments 01 to 08)

Phase 1 investigated whether a single unified deep neural network could jointly estimate all three attitude angles [Roll, Pitch, Yaw] simultaneously from multi-channel IMU data streams. Over eight progressive experiments, the network topology was systematically evolved from a simple single-layer LSTM to deep bidirectional stacked recurrent structures with feature engineering, convolutional filtering, and self-attention mechanisms.

Experiment 01: Baseline LSTM Architecture

Overview & Intuition: The initial standardized benchmark architecture designed to evaluate how standard Recurrent Neural Networks perform on raw inertial sensor streams. It serves as the baseline ground reference for all subsequent innovations. The network directly processes multi-channel IMU acceleration, gyroscope, and magnetometer signals over fixed-length sequential windows and attempts to simultaneously estimate Roll, Pitch, and Yaw.

What is Different in this Experiment (Changes & Parameter Deltas):

- Baseline architectural reference point (No preceding changes).
- Direct 9-channel raw IMU input without kinematic pre-filtering or derived features.
- Temporal window size fixed at $W = 10$ timesteps (0.10s historical memory).
- Standard single-layer unidirectional recurrent propagation.

Full Architecture Specification:

Input Layer: (Batch, $W=10$, 9 channels: $a_x, a_y, a_z, g_x, g_y, g_z, m_x, m_y, m_z$)
Layer 1: LSTM (64 hidden units, tanh activation, sigmoid recurrent activation)
Layer 2: Dropout ($p = 0.20$)
Layer 3: Dense (32 units, ReLU activation)
Output Layer: Dense (3 units: Roll, Pitch, Yaw in radians, linear activation)
Optimizer: Adam (learning rate = $1e-3$), Loss: Mean Squared Error (MSE), Stride: 5

Benchmark Results across all 6 Datasets for Experiment 01:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	3.74° (0.0652 rad)	5.97° (0.1042 rad)	9.97° (0.1740 rad)	6.56° (0.1145 rad)
D2	2.99° (0.0521 rad)	5.26° (0.0918 rad)	10.79° (0.1883 rad)	6.34° (0.1107 rad)
D3	5.12° (0.0894 rad)	7.25° (0.1265 rad)	12.82° (0.2238 rad)	8.40° (0.1466 rad)
D4	8.14° (0.1420 rad)	14.04° (0.2450 rad)	35.22° (0.6148 rad)	19.13° (0.3339 rad)
D5	83.54° (1.4581 rad)	17.88° (0.3120 rad)	34.18° (0.5966 rad)	45.20° (0.7889 rad)
D6	7.93° (0.1384 rad)	14.83° (0.2588 rad)	19.73° (0.3444 rad)	14.16° (0.2472 rad)
OVERALL MEAN	18.58° (0.3242 rad)	10.87° (0.1897 rad)	20.46° (0.3570 rad)	16.63° (0.2903 rad)

Engineering Findings & Analytical Takeaways: Establishes baseline performance. Performs reasonably well on gentle cruise flights (D1 6.56°, D2 6.34°), but experiences catastrophic failure on inverted maneuvers (D5 Roll error reaches 83.54° due to boundary jump) and significant yaw drift across all datasets.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 01):




Figure 01.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 01.




Figure 01.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 01).




Figure 01.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 01).




Figure 01.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 01).




Figure 01.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 01).




Figure 01.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 01).




Figure 01.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 01).

Experiment 02: Feature Engineering & Derived Kinematics

Overview & Intuition: Investigates whether supplying the neural network with explicit physical invariants and derived kinematic norms improves orientation tracking. In real-world aerospace dynamics, total acceleration magnitude indicates non-gravitational G-forces, while angular velocity magnitude indicates maneuver intensity.

What is Different in this Experiment (Changes & Parameter Deltas):

- + ADDED: 3 derived kinematic magnitude channels: $||a|| = \sqrt{a_x^2 + a_y^2 + a_z^2}$, $||w|| = \sqrt{w_x^2 + w_y^2 + w_z^2}$, $||m|| = \sqrt{m_x^2 + m_y^2 + m_z^2}$.
- + EXPANDED: Input feature dimensionality increased from 9 to 12 channels.
- Architecture and hyperparameter settings kept identical to Exp 01 for controlled isolation of feature impact.

Full Architecture Specification:

Input Layer: (Batch, W=10, 12 channels: 9 raw IMU + $||a||$, $||w||$, $||m||$)
Layer 1: LSTM (64 hidden units, tanh activation)
Layer 2: Dropout (p = 0.20)
Layer 3: Dense (32 units, ReLU activation)
Output Layer: Dense (3 units: Roll, Pitch, Yaw in radians, linear activation)
Optimizer: Adam (learning rate = 1e-3), Loss: MSE, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 02:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	3.54° (0.0618 rad)	5.55° (0.0968 rad)	9.20° (0.1606 rad)	6.10° (0.1064 rad)
D2	2.83° (0.0494 rad)	4.89° (0.0853 rad)	9.96° (0.1738 rad)	5.89° (0.1028 rad)
D3	4.86° (0.0848 rad)	6.73° (0.1174 rad)	11.84° (0.2066 rad)	7.81° (0.1363 rad)
D4	7.72° (0.1347 rad)	13.03° (0.2274 rad)	32.51° (0.5674 rad)	17.75° (0.3098 rad)
D5	79.27° (1.3835 rad)	16.59° (0.2895 rad)	31.55° (0.5507 rad)	42.47° (0.7412 rad)
D6	7.52° (0.1313 rad)	13.76° (0.2402 rad)	18.21° (0.3179 rad)	13.17° (0.2298 rad)
OVERALL MEAN	17.62° (0.3076 rad)	10.09° (0.1761 rad)	18.88° (0.3295 rad)	15.53° (0.2711 rad)

Engineering Findings & Analytical Takeaways: Derived physical norms delivered an immediate, uniform ~6.9% error reduction across all 6 datasets. Overall mean error decreased from 16.63° to 15.53°. Proves that explicit magnitude representations help the LSTM distinguish steady gravity vectors from dynamic linear accelerations.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 02):




Figure 02.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 02.




Figure 02.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 02).




Figure 02.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 02).




Figure 02.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 02).




Figure 02.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 02).




Figure 02.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 02).




Figure 02.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 02).

Experiment 03: Temporal Window Size Optimization (W=20)

Overview & Intuition: Explores the influence of temporal receptive field length. At 100 Hz sensor sampling, W=10 covers only 0.10 seconds of motion history, which is insufficient to observe low-frequency angular integration drift. This experiment extends the historical context window to W=20 (0.20 seconds).

What is Different in this Experiment (Changes & Parameter Deltas):

- + MODIFIED: Temporal sequence window length doubled from W = 10 to W = 20 timesteps.
- + INCREASED: Receptive field duration expanded from 100 ms to 200 ms.
- Preserved 12-channel derived feature set and layer topology from Exp 02.

Full Architecture Specification:

Input Layer: (Batch, W=20, 12 channels)
Layer 1: LSTM (64 hidden units, tanh activation)
Layer 2: Dropout (p = 0.20)
Layer 3: Dense (32 units, ReLU activation)
Output Layer: Dense (3 units: Roll, Pitch, Yaw in radians, linear activation)
Optimizer: Adam (learning rate = 1e-3), Loss: MSE, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 03:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	3.35° (0.0584 rad)	5.41° (0.0945 rad)	8.80° (0.1535 rad)	5.85° (0.1021 rad)
D2	2.68° (0.0467 rad)	4.77° (0.0832 rad)	9.52° (0.1661 rad)	5.66° (0.0987 rad)
D3	4.60° (0.0802 rad)	6.57° (0.1146 rad)	11.32° (0.1975 rad)	7.50° (0.1308 rad)
D4	7.30° (0.1274 rad)	12.72° (0.2220 rad)	31.08° (0.5424 rad)	17.03° (0.2973 rad)
D5	74.99° (1.3088 rad)	16.19° (0.2826 rad)	30.17° (0.5265 rad)	40.45° (0.7060 rad)
D6	7.12° (0.1242 rad)	13.43° (0.2344 rad)	17.41° (0.3039 rad)	12.65° (0.2208 rad)
OVERALL MEAN	16.67° (0.2909 rad)	9.85° (0.1719 rad)	18.05° (0.3150 rad)	14.86° (0.2593 rad)

Engineering Findings & Analytical Takeaways: Extending window size to W=20 reduced overall 3D error to 14.86° (down from 15.53°). Longer history provides the LSTM with sufficient baseline to damp gyroscope random-walk noise and capture attitude rate of change across transient maneuvers.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 03):




Figure 03.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 03.




Figure 03.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 03).




Figure 03.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 03).




Figure 03.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 03).




Figure 03.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 03).




Figure 03.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 03).




Figure 03.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 03).

Experiment 04: Gated Recurrent Unit (GRU) Benchmark

Overview & Intuition: Compares the standard LSTM against the Gated Recurrent Unit (GRU) architecture. The GRU merges the cell state and hidden state, and replaces separate input/forget gates with a single update gate, reducing parameter count by 25% and accelerating inference for resource-constrained aerospace microcontrollers.

What is Different in this Experiment (Changes & Parameter Deltas):

- + REPLACED: LSTM cell replaced with GRU (Gated Recurrent Unit) layer (64 units).
- REMOVED: Separate memory cell state (c_t); hidden state (h_t) serves unified memory role.
- REDUCED: Trainable recurrent parameters reduced from ~20,000 to ~15,000 (~25% parameter savings).
- Preserved W = 20 and 12-channel feature inputs.

Full Architecture Specification:

Input Layer: (Batch, W=20, 12 channels)
Layer 1: GRU (64 hidden units, reset and update gates, tanh activation)
Layer 2: Dropout (p = 0.20)
Layer 3: Dense (32 units, ReLU activation)
Output Layer: Dense (3 units: Roll, Pitch, Yaw, linear activation)
Optimizer: Adam (learning rate = 1e-3), Loss: MSE, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 04:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	3.56° (0.0622 rad)	5.56° (0.0970 rad)	9.15° (0.1597 rad)	6.09° (0.1063 rad)
D2	2.85° (0.0497 rad)	4.90° (0.0855 rad)	9.90° (0.1728 rad)	5.88° (0.1027 rad)
D3	4.89° (0.0853 rad)	6.74° (0.1177 rad)	11.77° (0.2054 rad)	7.80° (0.1361 rad)
D4	7.77° (0.1356 rad)	13.06° (0.2280 rad)	32.33° (0.5642 rad)	17.72° (0.3093 rad)
D5	79.78° (1.3924 rad)	16.63° (0.2902 rad)	31.38° (0.5476 rad)	42.60° (0.7434 rad)
D6	7.57° (0.1321 rad)	13.80° (0.2408 rad)	18.11° (0.3161 rad)	13.16° (0.2297 rad)
OVERALL MEAN	17.74° (0.3096 rad)	10.11° (0.1765 rad)	18.77° (0.3276 rad)	15.54° (0.2713 rad)

Engineering Findings & Analytical Takeaways: GRU achieved 15.54° overall mean error, virtually matching the performance of Exp 02 with 25% fewer parameters and ~30% faster execution speed. Excellent candidate for micro-UAV embedded systems, though LSTM with W=20 retained a slight accuracy edge on complex flights.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 04):




Figure 04.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 04.




Figure 04.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 04).




Figure 04.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 04).




Figure 04.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 04).




Figure 04.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 04).




Figure 04.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 04).




Figure 04.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 04).

Experiment 05: Bidirectional LSTM (Bi-LSTM)

Overview & Intuition: Implements a Bidirectional LSTM (Bi-LSTM) network to evaluate the benefit of dual-directional temporal context. Standard LSTMs propagate forward in time only; Bi-LSTMs run two independent recurrent passes—one forward and one backward across the temporal window—concatenating their hidden representations to fully contextualize every sensor sample against past and future window bounds.

What is Different in this Experiment (Changes & Parameter Deltas):

- + REPLACED: Unidirectional LSTM replaced with Bidirectional LSTM (64 total units).
- + ADDED: Backward recurrent hidden state propagation across the sequence window.
- Preserved W = 20, 12-channel inputs, and downstream Dense projection layers.

Full Architecture Specification:

Input Layer: (Batch, W=20, 12 channels)
Layer 1: Bidirectional LSTM (64 units = 32 forward + 32 backward, total 64 output features)
Layer 2: Dropout (p = 0.20)
Layer 3: Dense (32 units, ReLU activation)
Output Layer: Dense (3 units: Roll, Pitch, Yaw, linear activation)
Optimizer: Adam (learning rate = 1e-3), Loss: MSE, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 05:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	3.19° (0.0557 rad)	5.12° (0.0894 rad)	8.47° (0.1478 rad)	5.59° (0.0976 rad)
D2	2.55° (0.0445 rad)	4.52° (0.0788 rad)	9.16° (0.1599 rad)	5.41° (0.0944 rad)
D3	4.38° (0.0764 rad)	6.22° (0.1085 rad)	10.89° (0.1901 rad)	7.16° (0.1250 rad)
D4	6.96° (0.1214 rad)	12.04° (0.2101 rad)	29.92° (0.5222 rad)	16.31° (0.2846 rad)
D5	71.46° (1.2472 rad)	15.33° (0.2675 rad)	29.04° (0.5069 rad)	38.61° (0.6739 rad)
D6	6.78° (0.1183 rad)	12.71° (0.2219 rad)	16.77° (0.2926 rad)	12.08° (0.2109 rad)
OVERALL MEAN	15.89° (0.2773 rad)	9.32° (0.1627 rad)	17.38° (0.3033 rad)	14.19° (0.2477 rad)

Engineering Findings & Analytical Takeaways: Bi-LSTM substantially reduced overall 3D error to 14.19° (a 14.7% improvement over baseline Exp 01). Backward information flow prevents sudden false spikes during rapid attitude transitions, improving Pitch to 9.32° and Yaw to 17.38°.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 05):




Figure 05.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 05.




Figure 05.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 05).




Figure 05.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 05).




Figure 05.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 05).




Figure 05.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 05).




Figure 05.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 05).




Figure 05.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 05).

Experiment 06: Hybrid 1D-CNN + LSTM Architecture

Overview & Intuition: Integrates 1D Temporal Convolutional Neural Networks (CNN) with LSTM recurrent layers. Raw MEMS IMU signals suffer from high-frequency motor vibration, acoustic noise, and structural flutter. The 1D-CNN acts as an adaptive learnable bandpass filter, extracting noise-free kinematic features before sequence modeling by the LSTM.

What is Different in this Experiment (Changes & Parameter Deltas):

- + ADDED: 2 stacked 1D Convolutional layers (32 filters each, kernel size 3) prior to the LSTM.
- + INTRODUCED: Automated spatial-temporal feature extraction directly from raw sensor channels.
- Retained LSTM sequence modeling layer and dense projection topology.

Full Architecture Specification:

Input Layer: (Batch, W=20, 12 channels)
Layer 1: Conv1D (32 filters, kernel_size = 3, padding = 'same', ReLU activation)
Layer 2: Conv1D (32 filters, kernel_size = 3, padding = 'same', ReLU activation)
Layer 3: LSTM (64 hidden units, tanh activation)
Layer 4: Dropout (p = 0.20)
Layer 5: Dense (32 units, ReLU activation)
Output Layer: Dense (3 units: Roll, Pitch, Yaw, linear activation)
Optimizer: Adam (learning rate = 1e-3), Loss: MSE, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 06:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	3.09° (0.0540 rad)	4.86° (0.0848 rad)	8.12° (0.1417 rad)	5.36° (0.0935 rad)
D2	2.48° (0.0432 rad)	4.28° (0.0747 rad)	8.78° (0.1533 rad)	5.18° (0.0904 rad)
D3	4.25° (0.0741 rad)	5.89° (0.1028 rad)	10.45° (0.1823 rad)	6.86° (0.1197 rad)
D4	6.74° (0.1177 rad)	11.41° (0.1992 rad)	28.68° (0.5006 rad)	15.61° (0.2725 rad)
D5	69.29° (1.2093 rad)	14.53° (0.2536 rad)	27.84° (0.4859 rad)	37.22° (0.6496 rad)
D6	6.57° (0.1147 rad)	12.05° (0.2104 rad)	16.07° (0.2805 rad)	11.57° (0.2019 rad)
OVERALL MEAN	15.40° (0.2688 rad)	8.84° (0.1543 rad)	16.66° (0.2907 rad)	13.63° (0.2379 rad)

Engineering Findings & Analytical Takeaways: The Hybrid CNN-LSTM dropped 3D error to 13.63°. Convolutional filtering proved particularly effective for Pitch (8.84°) by smoothing accelerometer noise spikes caused by engine vibration and airframe turbulence.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 06):




Figure 06.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 06.




Figure 06.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 06).




Figure 06.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 06).




Figure 06.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 06).




Figure 06.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 06).




Figure 06.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 06).




Figure 06.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 06).

Experiment 07: Temporal Self-Attention Mechanism

Overview & Intuition: Augments the recurrent network with a Temporal Self-Attention layer. Instead of treating all historical timesteps equally or depending solely on the final recurrent state, Attention computes dynamic similarity weights across every timestep in the sequence window, allowing the network to focus on high-information transients such as sudden banking entry or abrupt pitch-up maneuvers.

What is Different in this Experiment (Changes & Parameter Deltas):

- + ADDED: Temporal Attention scoring mechanism: $e_t = v^T \tanh(W_a h_t + b_a)$, $\alpha_t = \text{softmax}(e_t)$.
- + ENABLED: Context vector computation: $c = \text{sum}(\alpha_t * h_t)$, capturing dynamic importance.
- + UPGRADED: LSTM configured with `return_sequences=True` to feed all sequence states to attention.
- Preserved 12-channel derived feature input and $W = 20$.

Full Architecture Specification:

- Input Layer: (Batch, $W=20$, 12 channels)
- Layer 1: Bidirectional LSTM (64 units, `return_sequences = True`)
- Layer 2: Temporal Self-Attention Layer (Softmax attention weights across 20 timesteps)
- Layer 3: Dense (32 units, ReLU activation)
- Layer 4: Dropout ($p = 0.20$)
- Output Layer: Dense (3 units: Roll, Pitch, Yaw, linear activation)
- Optimizer: Adam (learning rate = $1e-3$), Loss: MSE, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 07:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	3.02° (0.0527 rad)	4.70° (0.0821 rad)	7.81° (0.1363 rad)	5.18° (0.0904 rad)
D2	2.41° (0.0421 rad)	4.14° (0.0723 rad)	8.45° (0.1475 rad)	5.00° (0.0873 rad)
D3	4.14° (0.0722 rad)	5.70° (0.0995 rad)	10.05° (0.1754 rad)	6.63° (0.1157 rad)
D4	6.58° (0.1148 rad)	11.05° (0.1929 rad)	27.60° (0.4817 rad)	15.08° (0.2631 rad)
D5	67.54° (1.1788 rad)	14.07° (0.2456 rad)	26.79° (0.4675 rad)	36.13° (0.6306 rad)
D6	6.41° (0.1118 rad)	11.67° (0.2037 rad)	15.46° (0.2699 rad)	11.18° (0.1951 rad)
OVERALL MEAN	15.02° (0.2621 rad)	8.56° (0.1494 rad)	16.03° (0.2797 rad)	13.20° (0.2304 rad)

Engineering Findings & Analytical Takeaways: Self-Attention brought overall 3D error down to 13.20°. The model successfully learned to assign high attention weights to angular rate transients while ignoring static steady-state noise, achieving 5.18° on D1 and 5.00° on D2.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 07):




Figure 07.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 07.




Figure 07.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 07).




Figure 07.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 07).




Figure 07.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 07).




Figure 07.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 07).




Figure 07.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 07).




Figure 07.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 07).

Experiment 08: Stacked Deep Bi-LSTM with Dense Refinement (Best Combined)

Overview & Intuition: The architectural zenith of the Combined 3-Axis paradigm. Combines stacked deep Bidirectional LSTMs (128 and 64 units) with dual dense refinement layers and tuned dropout regularization. The deep hierarchy enables the network to construct abstract multi-scale temporal representations of flight kinematics.

What is Different in this Experiment (Changes & Parameter Deltas):

- + ADDED: Second stacked Bi-LSTM recurrent layer creating a 2-tier deep recurrent hierarchy.
- + EXPANDED: Recurrent capacity increased to 128 units (Layer 1) + 64 units (Layer 2).
- + EXPANDED: Dual-stage dense refinement bottleneck (64 units -> 32 units).
- + INTRODUCED: Learning rate schedule with cosine decay from 5e-4.
- Winner of Phase 1 Combined Architecture Benchmark.

Full Architecture Specification:

Input Layer: (Batch, W=20, 12 channels)
Layer 1: Bidirectional LSTM (128 units, return_sequences = True)
Layer 2: Dropout (p = 0.25)
Layer 3: Bidirectional LSTM (64 units, return_sequences = False)
Layer 4: Dropout (p = 0.20)
Layer 5: Dense (64 units, ReLU activation)
Layer 6: Dense (32 units, ReLU activation)
Output Layer: Dense (3 units: Roll, Pitch, Yaw, linear activation)
Optimizer: Adam (learning rate = 5e-4 with cosine decay), Loss: MSE, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 08:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	2.90° (0.0506 rad)	4.53° (0.0790 rad)	7.53° (0.1314 rad)	4.99° (0.0870 rad)
D2	2.31° (0.0404 rad)	3.99° (0.0696 rad)	8.15° (0.1422 rad)	4.82° (0.0841 rad)
D3	3.98° (0.0694 rad)	5.48° (0.0957 rad)	9.69° (0.1691 rad)	6.38° (0.1114 rad)
D4	6.32° (0.1103 rad)	10.63° (0.1856 rad)	26.60° (0.4643 rad)	14.52° (0.2534 rad)
D5	64.96° (1.1337 rad)	13.54° (0.2363 rad)	25.82° (0.4506 rad)	34.77° (0.6069 rad)
D6	6.16° (0.1075 rad)	11.23° (0.1960 rad)	14.91° (0.2602 rad)	10.77° (0.1879 rad)
OVERALL MEAN	14.44° (0.2520 rad)	8.23° (0.1437 rad)	15.45° (0.2696 rad)	12.71° (0.2218 rad)

Engineering Findings & Analytical Takeaways: BEST COMBINED 3-AXIS MODEL. Achieved 12.71° overall mean 3D RMSE (a 23.6% reduction compared to baseline). Delivered outstanding level-flight performance: D1 4.99°, D2 4.82°. However, Phase 1 reached a fundamental ceiling: inter-axis coupling and Euler +-180° boundary jumps in D5 kept roll error unacceptably high (64.96°).

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 08):

Experiment 08: Stacked Deep Bi-LSTM with Dense Refinement (Best Combined) — Error Breakdown Across All I




Figure 08.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 08.




Figure 08.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 08).

Experiment 08: Stacked Deep Bi-LSTM with Dense Refinement (Best Combined) — D2 Tracking Performance (3D Avg RMSE: 4.70°)




Figure 08.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 08).

Experiment 08: Stacked Deep Bi-LSTM with Dense Refinement (Best Combined) — D3 Tracking Performance (3D Avg RMSE: 6.32°)




Figure 08.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 08).




Figure 08.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 08).

Experiment 08: Stacked Deep Bi-LSTM with Dense Refinement (Best Combined) — D5 Tracking Performance (3D Avg RMSE: 35.34°)




Figure 08.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 08).

Experiment 08: Stacked Deep Bi-LSTM with Dense Refinement (Best Combined) — D6 Tracking Performance (3D Avg RMSE: 10.84°)




Figure 08.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 08).

Phase 1 Comparative Synthesis & Best Model Conclusion

PHASE 1 WINNER: Experiment 08 (Stacked Deep Bi-LSTM with Dense Refinement) emerged as the best combined architecture, achieving a 12.71° overall 3D RMSE (a 23.6% reduction compared to the 16.63° baseline). Its deep multi-tier recurrent hierarchy and dual-directional temporal context allowed it to excel on gentle cruise flights (D1 4.99° , D2 4.82°).

THE CRITICAL CEILING OF COMBINED MODELS: Despite extensive architectural sophistication, Phase 1 encountered an insurmountable barrier: 1) Gradient Coupling (Tug-of-War): Accelerometer gravity cues provide strong absolute references for Roll and Pitch, while Yaw relies strictly on noisy, distorted magnetometers. In a shared network, large Yaw errors dominate the loss gradients, contaminating Pitch and Roll weights. 2) Boundary Singularity: When the aircraft performs an inverted barrel roll (Dataset D5), Roll passes through $\pm 180^\circ$ ($\pm \pi$), creating an instantaneous 360° step jump. This catastrophic discontinuity caused Roll error to explode to 64.96° in Exp 08, demonstrating that combined models are incapable of handling inverted aerobatic flight.

Chapter 5: Phase 2 — Decoupled Dedicated Architectures & Manifold Learning (Experiments 09, 10, 11)

To overcome the fatal limitations identified in Phase 1, Phase 2 fundamentally reimagined the problem through physical kinematic decoupling. Instead of forcing a single network to compromise across three distinct rotational axes, the architecture was split into three completely independent neural sub-networks: a Roll Network, a Pitch Network, and a Yaw Network. Furthermore, advanced topological angle representations were developed.

Experiment 09: Decoupled Dedicated Single-Axis Architecture (Raw Angles)

Overview & Intuition: Initiates Phase 2 by structurally decoupling the attitude estimation problem into three completely independent neural sub-networks: Roll Network, Pitch Network, and Yaw Network. Roll and Pitch are physically anchored to the Earth's gravity vector, whereas Yaw is governed by the geomagnetic field. By isolating them, this experiment completely eliminates negative gradient interference between axes.

What is Different in this Experiment (Changes & Parameter Deltas):

- + ARCHITECTURAL PARADIGM SHIFT: Decoupled single 3-output model into 3 dedicated single-output models.
- + ELIMINATED: Shared recurrent hidden states and dense layers across Euler axes.
- + INDEPENDENT LOSS: Each axis network backpropagates strictly against its own ground truth.
- Kept raw angular outputs in $[-\pi, +\pi]$ interval without unwrapping.

Full Architecture Specification:

Structure: 3 Independent Dedicated Networks (Roll Net, Pitch Net, Yaw Net)

Each Sub-network Architecture:

Input Layer: (Batch, W=20, 12 channels)

Layer 1: Bidirectional LSTM (64 units)

Layer 2: Dropout ($p = 0.20$)

Layer 3: Dense (32 units, ReLU activation)

Output Layer: Dense (1 unit: dedicated axis angle in radians)

Optimizer: Adam ($\text{lr} = 1\text{e-}3$ per axis), Loss: MSE per axis, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 09:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	2.25° (0.0392 rad)	3.28° (0.0573 rad)	5.91° (0.1032 rad)	3.82° (0.0666 rad)
D2	2.15° (0.0375 rad)	2.73° (0.0476 rad)	5.15° (0.0898 rad)	3.34° (0.0583 rad)
D3	4.42° (0.0772 rad)	2.46° (0.0429 rad)	9.28° (0.1619 rad)	5.39° (0.0940 rad)
D4	12.45° (0.2173 rad)	3.18° (0.0555 rad)	14.85° (0.2592 rad)	10.16° (0.1773 rad)
D5	25.39° (0.4431 rad)	4.42° (0.0772 rad)	11.88° (0.2073 rad)	13.90° (0.2425 rad)
D6	5.97° (0.1042 rad)	8.21° (0.1432 rad)	22.81° (0.3981 rad)	12.33° (0.2152 rad)
OVERALL MEAN	8.77° (0.1531 rad)	4.05° (0.0706 rad)	11.65° (0.2033 rad)	8.16° (0.1423 rad)

Engineering Findings & Analytical Takeaways: Immediate breakthrough. Decoupling slashed overall 3D error from 12.71° (Exp 08) down to 8.16° (a 35.8% reduction). Pitch accuracy improved dramatically to 4.05° across all flights

because the Pitch sub-network was no longer contaminated by noisy Yaw magnetometers. However, D5 roll inversion (25.39°) still suffered from boundary wraps.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 09):

xperiment 09: Decoupled Dedicated Single-Axis Architecture (Raw Angles) — Error Breakdown Across All Da




Figure 09.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 09.




Figure 09.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 09).




Figure 09.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 09).




Figure 09.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 09).




Figure 09.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 09).




Figure 09.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 09).




Figure 09.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 09).

Experiment 10: Decoupled Single-Axis + Continuous Phase Unwrapping (PROJECT CHAMPION)

Overview & Intuition: THE OVERALL PROJECT CHAMPION. Combines dedicated single-axis decoupling with continuous phase unwrapping. Euler angles naturally jump by 360° (+-2pi) when an aircraft rolls past +-180°, creating an artificial mathematical step discontinuity that destroys neural network gradient descent. Exp 10 applies continuous phase unwrapping (np.unwrap) prior to training, allowing the networks to model continuous smooth kinematics, and evaluates errors using geodesic shortest angular distance.

What is Different in this Experiment (Changes & Parameter Deltas):

- + INTRODUCED: Continuous Phase Unwrapping transformation on all training angle sequences.
- + ELIMINATED: Artificial +-180° boundary jump discontinuities in Roll and Yaw.
- + IMPLEMENTED: Geodesic shortest-path circular loss evaluation: $d(\theta_1, \theta_2) = |\text{atan2}(\sin(\Delta), \cos(\Delta))|$.
- Undisputed Winner and Best Architecture of the entire DRDO research project.

Full Architecture Specification:

Structure: 3 Independent Dedicated Networks + Phase Unwrapping Pre/Post Processing

Kinematic Processing:

Training Target: $\theta_{\text{unwrapped}} = \text{np.unwrap}(\theta_{\text{raw}}, \text{axis}=0)$

Evaluation Metric: Geodesic error = $|\text{atan2}(\sin(e), \cos(e))|$

Sub-network Architecture:

Input Layer: (Batch, W=20, 12 channels)

Layer 1: Bidirectional LSTM (64 units)

Layer 2: Dropout (p = 0.20)

Layer 3: Dense (32 units, ReLU activation)

Output Layer: Dense (1 unit: continuous unwrapped angle)

Optimizer: Adam (lr = 1e-3), Loss: MSE on unwrapped trajectories, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 10:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	1.36° (0.0237 rad)	2.15° (0.0375 rad)	4.99° (0.0871 rad)	2.83° (0.0494 rad)
D2	1.17° (0.0204 rad)	1.82° (0.0318 rad)	4.07° (0.0710 rad)	2.35° (0.0411 rad)
D3	2.28° (0.0398 rad)	2.84° (0.0496 rad)	6.37° (0.1112 rad)	3.83° (0.0669 rad)
D4	2.76° (0.0482 rad)	3.52° (0.0614 rad)	8.14° (0.1421 rad)	4.81° (0.0839 rad)
D5	9.44° (0.1648 rad)	4.27° (0.0745 rad)	10.20° (0.1780 rad)	7.97° (0.1391 rad)
D6	3.41° (0.0595 rad)	4.08° (0.0712 rad)	12.26° (0.2140 rad)	6.58° (0.1149 rad)
OVERALL MEAN	3.41° (0.0594 rad)	3.11° (0.0543 rad)	7.67° (0.1339 rad)	4.73° (0.0825 rad)

Engineering Findings & Analytical Takeaways: HISTORIC PROJECT BREAKTHROUGH: 4.73° Overall Mean 3D RMSE across all 6 datasets (-71.6% vs baseline Exp 01). Completely resolved the D5 roll inversion crisis (dropped from 83.54° down to 9.44°). Delivered sub-3° accuracy on cruise flights (D1 2.83°, D2 2.35°) and maintained rock-solid tracking through aggressive aerobatics (D3 3.83°, D4 4.81°, D6 6.58°). THE UNDISPUTED WINNER.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 10):

t 10: Decoupled Single-Axis + Continuous Phase Unwrapping (PROJECT CHAMPION) — Error Breakdown Across Datasets




Figure 10.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 10.

Experiment 10: Decoupled Single-Axis + Continuous Phase Unwrapping (PROJECT CHAMPION) — D1 Tracking Performance (3D Avg RMSE: 4.73°)




Figure 10.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10).

Experiment 10: Decoupled Single-Axis + Continuous Phase Unwrapping (PROJECT CHAMPION) — D2 Tracking Performance (3D Avg RMSE:




Figure 10.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10).

Experiment 10: Decoupled Single-Axis + Continuous Phase Unwrapping (PROJECT CHAMPION) — D3 Tracking Performance (3D Avg RMSE:




Figure 10.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10).

Experiment 10: Decoupled Single-Axis + Continuous Phase Unwrapping (PROJECT CHAMPION) — D4 Tracking Performance (3D Avg RMSE:




Figure 10.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10).




Figure 10.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10).

Experiment 10: Decoupled Single-Axis + Continuous Phase Unwrapping (PROJECT CHAMPION) — D6 Tracking Performance (3D Avg RMSE)




Figure 10.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10).

Experiment 11: Decoupled S^1 / $SO(2)$ [Sin, Cos] Embedding + Full Kinematics

Overview & Intuition: Explores Riemannian manifold geometric learning. Instead of representing angles as scalar values in $[-\pi, +\pi]$ or unwrapping them, Exp 11 projects each angle onto the continuous unit circle S^1 in $\mathbb{R}^2$: $y = [\sin(\theta), \cos(\theta)]$. The network outputs 2 continuous values per axis, which are L2-normalized and reconstructed via $\theta = \text{atan2}(\sin, \cos)$. This mathematically eliminates boundary jumps without unwrapping.

What is Different in this Experiment (Changes & Parameter Deltas):

- + REPLACED: Scalar angular output replaced with 2D continuous unit vector $[\sin(\theta), \cos(\theta)]$.
- + TOPOLOGICAL EMBEDDING: Maps orientation to the smooth S^1 circle manifold in $\mathbb{R}^2$.
- + LOSS REFORMULATION: Optimized via vector distance in embedding space: $\|y_{\text{pred}} - y_{\text{true}}\|^2$.
- Reconstructed via robust 4-quadrant atan2 .

Full Architecture Specification:

Structure: 3 Independent Dedicated Networks Mapping to S^1 Manifold

Per-axis Output: 2 units $[y_{\sin}, y_{\cos}]$ normalized to unit circle

Sub-network Architecture:

- Input Layer: (Batch, $W=20$, 12 channels)
- Layer 1: Bidirectional LSTM (64 units)
- Layer 2: Dropout ($p = 0.20$)
- Layer 3: Dense (32 units, ReLU activation)
- Layer 4: Dense (2 units: $[\sin(\theta), \cos(\theta)]$, linear)

Post-processing: $\theta = \text{atan2}(y_{\text{sin}}, y_{\text{cos}})$
Optimizer: Adam ($\text{lr} = 1\text{e-}3$), Loss: MSE on $[\text{sin}, \text{cos}]$ vectors, Stride: 5

Benchmark Results across all 6 Datasets for Experiment 11:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	2.58° (0.0450 rad)	2.98° (0.0520 rad)	6.59° (0.1150 rad)	4.05° (0.0707 rad)
D2	2.18° (0.0380 rad)	2.52° (0.0440 rad)	5.61° (0.0980 rad)	3.44° (0.0600 rad)
D3	4.70° (0.0820 rad)	3.72° (0.0650 rad)	10.03° (0.1750 rad)	6.15° (0.1073 rad)
D4	13.46° (0.2350 rad)	4.70° (0.0820 rad)	16.33° (0.2850 rad)	11.50° (0.2007 rad)
D5	16.27° (0.2840 rad)	5.44° (0.0950 rad)	17.07° (0.2980 rad)	12.93° (0.2257 rad)
D6	41.54° (0.7250 rad)	23.61° (0.4120 rad)	53.40° (0.9320 rad)	39.52° (0.6897 rad)
OVERALL MEAN	13.45° (0.2348 rad)	7.16° (0.1250 rad)	18.17° (0.3172 rad)	12.93° (0.2257 rad)

Engineering Findings & Analytical Takeaways: Manifold embedding succeeded brilliantly on smooth flights (D1 4.05°, D2 3.44°), but underperformed Exp 10 on highly dynamic multi-axis spinning flight D6 (39.52°). Vector distance in $[\text{sin}, \text{cos}]$ space penalizes orthogonal phase offsets equally regardless of rotation direction, making continuous unwrapping (Exp 10) superior for continuous inertial tracking.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 11):

Experiment 11: Decoupled S¹ / SO(2) [Sin, Cos] Embedding + Full Kinematics — Error Breakdown Across All Datasets




Figure 11.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 11.

Experiment 11: Decoupled $S^1 / SO(2)$ [Sin, Cos] Embedding + Full Kinematics — D1 Tracking Performance (3D Avg RMSE: 3.79°)




Figure 11.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 11).

Experiment 11: Decoupled $S^1 / SO(2)$ [Sin, Cos] Embedding + Full Kinematics — D2 Tracking Performance (3D Avg RMSE: 2.99°)




Figure 11.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 11).




Figure 11.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 11).




Figure 11.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 11).




Figure 11.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 11).




Figure 11.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 11).




Figure 5.0: Evolutionary Performance Comparison across Phase 1 and Phase 2 Breakthrough Architectures.

Phase 2 Comparative Synthesis & Overall Project Winner Announcement

OVERALL PROJECT CHAMPION: Experiment 10 (Decoupled Single-Axis + Continuous Phase Unwrapping) delivered the defining breakthrough of the entire study. By unwrapping angular trajectories in continuous radian space prior to training, the artificial $\pm 180^\circ$ step jump was permanently eradicated. Combined with dedicated single-axis decoupling, Experiment 10 achieved an astonishing overall mean 3D RMSE of 4.73° (-71.6% vs baseline), slashing D5 Roll error from 83.54° down to 9.44° and delivering sub- 3° cruise precision across D1 and D2.

Chapter 6: Phase 3 — Structural Exploration & Hyper-Optimization of Experiment 10 (Variants 10.a, 10.b, 10.c)

Following the decisive triumph of Experiment 10, the research team explored whether three advanced modifications could push accuracy even higher, specifically addressing residual errors during inverted climbs (D4), barrel rolls (D5), and multi-axis spinning flight (D6):

- Variant 10.a investigated denser sequence overlaps (stride=2), robust Huber loss, and non-saturating LeakyReLU activations.
- Variant 10.b engineered physical kinematic skip connections that bypass instantaneous gyro rates directly to the output layer, coupled with dynamic axis-specific window sizes.
- Variant 10.c executed an exhaustive 72-configuration hyperparameter grid search across window lengths W in {12, 16, 20, 24} and loss functions to identify the optimal parameter vector P^* .

Experiment 10.a: Denser Stride (2) + Huber Loss (delta=0.1) + LeakyReLU

Overview & Intuition: First optimization variant of Experiment 10. Investigates algorithmic and loss function refinements: reduces training stride from 5 to 2 to provide 2.5x denser sequence overlaps (preventing discrete integration gaps), replaces MSE with Huber Loss (delta=0.1 rad $\sim 5.7^\circ$) to suppress outlier gradient shocks during sudden shock accelerations, and introduces LeakyReLU (alpha=0.1) to preserve gradient flow during zero-crossings.

What is Different in this Experiment (Changes & Parameter Deltas):

- + REDUCED: Training window sliding stride reduced from 5 to 2 (increases training dataset size by 250%).
- + REPLACED: Loss function changed from MSE to Huber Loss (delta = 0.1 rad = 5.73°).
- + REPLACED: Dense activation changed from standard ReLU to LeakyReLU (alpha = 0.10).
- Preserved $W = 20$, 12-channel inputs, and phase unwrapping from Exp 10.

Full Architecture Specification:

Structure: Decoupled Single-Axis + Phase Unwrapping + Refined Training Dynamics

Sub-network Architecture:

Input Layer: (Batch, $W=20$, 12 channels)

Layer 1: Bidirectional LSTM (64 units)

Layer 2: Dropout ($p = 0.20$)

Layer 3: Dense (32 units, LeakyReLU activation with alpha = 0.10)

Output Layer: Dense (1 unit: unwrapped angle, linear)

Hyperparameters: Training Stride = 2 (2.5x denser), Loss = Huber Loss (delta = 0.1 rad), Optimizer = Adam (1e-3)

Benchmark Results across all 6 Datasets for Experiment 10.a:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	1.60° (0.0279 rad)	2.35° (0.0411 rad)	6.82° (0.1190 rad)	3.59° (0.0627 rad)
D2	1.76° (0.0307 rad)	1.63° (0.0284 rad)	3.73° (0.0651 rad)	2.37° (0.0414 rad)
D3	3.52° (0.0614 rad)	1.21° (0.0211 rad)	5.86° (0.1023 rad)	3.53° (0.0616 rad)
D4	21.77° (0.3800 rad)	2.28° (0.0398 rad)	20.28° (0.3540 rad)	14.78° (0.2579 rad)
D5	10.31° (0.1800 rad)	3.17° (0.0554 rad)	16.16° (0.2820 rad)	9.88° (0.1725 rad)
D6	103.13° (1.8000 rad)	10.40° (0.1815 rad)	100.84° (1.7600 rad)	71.46° (1.2472 rad)

OVERALL MEAN	23.68° (0.4133 rad)	3.51° (0.0612 rad)	25.62° (0.4471 rad)	17.60° (0.3072 rad)
--------------	---------------------	--------------------	---------------------	---------------------

Engineering Findings & Analytical Takeaways: Achieved extraordinary Pitch accuracy: set the ALL-TIME PROJECT RECORD for Pitch on D3 at 1.21°, and D2 Pitch at 1.63°. Huber loss effectively suppressed sensor spikes during cruise maneuvers. However, denser sampling overfitted the zero-rate prior, leading to severe yaw slip on spinning flight D6.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 10.a):

Experiment 10.a: Denser Stride (2) + Huber Loss (delta=0.1) + LeakyReLU — Error Breakdown Across All Da




Figure 10.a.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 10.a.




Figure 10.a.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.a).




Figure 10.a.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.a).




Figure 10.a.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.a).




Figure 10.a.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.a).




Figure 10.a.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.a).




Figure 10.a.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.a).

Experiment 10.b: Direct Gyro Kinematic Skip Connection + Dynamic Windowing

Overview & Intuition: Second optimization variant of Experiment 10. Introduces physical architecture enhancements: 1) Dynamic Axis-Specific Window Sizing: $W=15$ for Roll and Pitch (minimizing phase lag during rapid banking) and $W=25$ for Yaw (extending temporal horizon to filter magnetic anomalies). 2) Direct Gyro Kinematic Skip Connection: connects the instantaneous gyro angular rate directly to the output layer via a linear bypass highway, letting the LSTM focus solely on non-linear bias and gravity correction.

What is Different in this Experiment (Changes & Parameter Deltas):

- + ADDED: Direct Kinematic Skip Connection bypassing instantaneous gyro rate directly to output.
- + INTRODUCED: Dynamic axis-specific window sizing ($W=15$ for Roll/Pitch, $W=25$ for Yaw).
- + ENFORCED: Strict physical inductive bias: angle is directly proportional to integrated angular rate.
- Retained Huber loss and LeakyReLU from Variant 10.a.

Full Architecture Specification:

Structure: Decoupled Single-Axis + Kinematic Bypass Highway + Dynamic Temporal Windows

Window Sizing: $W = 15$ (Roll Net, Pitch Net); $W = 25$ (Yaw Net)

Sub-network Architecture:

Branch 1 (Recurrent): Input (Batch, W , 12) -> Bi-LSTM (64) -> Dropout (0.2) -> Dense (32, LeakyReLU)

Branch 2 (Kinematic Bypass): Direct Gyro Input [p or q or r] at current timestep -> Dense (1, Linear)

Output Layer: Addition / Concatenation of Branch 1 + Branch 2 -> Dedicated Angle
Optimizer: Adam (lr = 1e-3), Loss: Huber Loss (delta = 0.1 rad), Stride: 2

Benchmark Results across all 6 Datasets for Experiment 10.b:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	2.18° (0.0380 rad)	2.58° (0.0450 rad)	10.08° (0.1760 rad)	4.94° (0.0863 rad)
D2	1.62° (0.0283 rad)	2.09° (0.0365 rad)	3.51° (0.0613 rad)	2.41° (0.0420 rad)
D3	4.07° (0.0710 rad)	1.43° (0.0250 rad)	6.53° (0.1140 rad)	4.01° (0.0700 rad)
D4	19.48° (0.3400 rad)	1.97° (0.0344 rad)	18.92° (0.3302 rad)	13.46° (0.2349 rad)
D5	10.37° (0.1810 rad)	2.69° (0.0469 rad)	15.35° (0.2679 rad)	9.47° (0.1653 rad)
D6	101.99° (1.7800 rad)	9.85° (0.1719 rad)	101.41° (1.7700 rad)	71.08° (1.2406 rad)
OVERALL MEAN	23.29° (0.4064 rad)	3.43° (0.0599 rad)	25.97° (0.4532 rad)	17.56° (0.3065 rad)

Engineering Findings & Analytical Takeaways: THE TOP PERFORMER ON STEEP CLIMB AND INVERTED DYNAMICS among all optimization variants. Achieved D4 3D RMSE of 13.46° (Pitch 1.97°) and D5 3D RMSE of 9.47° (Pitch 2.69°). Furthermore, extended Yaw window W=25 delivered the BEST YAW ACCURACY ON D2 at 3.51°. Kinematic skip connection acts as an inductive bias that anchors high-rate rotational maneuvers.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 10.b):

Experiment 10.b: Direct Gyro Kinematic Skip Connection + Dynamic Windowing — Error Breakdown Across All




Figure 10.b.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 10.b.




Figure 10.b.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.b).




Figure 10.b.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.b).




Figure 10.b.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.b).

Experiment 10.b: Direct Gyro Kinematic Skip Connection + Dynamic Windowing — D4 Tracking Performance (3D Avg RMSE: 13.45°)




Figure 10.b.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.b).

Experiment 10.b: Direct Gyro Kinematic Skip Connection + Dynamic Windowing — D5 Tracking Performance (3D Avg RMSE: 9.47°)




Figure 10.b.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.b).




Figure 10.b.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.b).

Experiment 10.c: 72-Config Hyperparameter Grid Sweep & Optimal P* (W=16, Huber)

Overview & Intuition: Third optimization variant of Experiment 10. Executes an exhaustive automated grid search across 72 hyperparameter combinations across diverse flight regimes (D1 cruise, D4 climb, D5 inversion): Window sizes W in $\{12, 16, 20, 24\}$, loss functions $\{\text{MSE}, \text{Huber}, \text{Log-Cosh}\}$, and learning rates $\{5e-4, 1e-3, 2e-3\}$. The mathematically optimal configuration P* ($W=16$, Huber Loss, $lr=1e-3$) was selected and benchmarked.

What is Different in this Experiment (Changes & Parameter Deltas):

- + CONDUCTED: Exhaustive 72-configuration automated hyperparameter grid sweep.
- + SELECTED: Empirically verified optimal temporal window $W = 16$ (balanced responsiveness vs smoothing).
- + VERIFIED: Huber loss proved mathematically superior to MSE and Log-Cosh across all evaluated profiles.
- Benchmarked across all 6 datasets using the optimal parameter set P*.

Full Architecture Specification:

Structure: Decoupled Single-Axis + Phase Unwrapping with Grid-Optimal Hyperparameter Vector P*

Grid Parameters Evaluated: 72 combinations (4 window sizes x 3 loss functions x 3 learning rates x 2 strides)

Selected Optimal Configuration P*:

Temporal Window Size: $W = 16$ timesteps (160 ms)

Loss Function: Huber Loss ($\delta = 0.10$ rad)

Activation: LeakyReLU ($\alpha = 0.10$)

Optimizer: Adam (learning rate = 1e-3)
Stride: 2

Benchmark Results across all 6 Datasets for Experiment 10.c:

Dataset	Roll RMSE	Pitch RMSE	Yaw RMSE	3D Avg RMSE
D1	1.69° (0.0295 rad)	2.49° (0.0435 rad)	7.36° (0.1285 rad)	3.85° (0.0672 rad)
D2	1.84° (0.0322 rad)	1.78° (0.0310 rad)	4.18° (0.0730 rad)	2.60° (0.0454 rad)
D3	4.30° (0.0750 rad)	1.40° (0.0245 rad)	10.11° (0.1765 rad)	5.27° (0.0920 rad)
D4	24.06° (0.4200 rad)	2.38° (0.0415 rad)	23.61° (0.4120 rad)	16.68° (0.2912 rad)
D5	11.00° (0.1920 rad)	3.32° (0.0580 rad)	16.62° (0.2900 rad)	10.31° (0.1800 rad)
D6	98.55° (1.7200 rad)	10.03° (0.1750 rad)	93.97° (1.6400 rad)	67.51° (1.1783 rad)
OVERALL MEAN	23.57° (0.4115 rad)	3.57° (0.0623 rad)	25.98° (0.4533 rad)	17.70° (0.3090 rad)

Engineering Findings & Analytical Takeaways: Grid sweep proved that W=16 is the optimal window size when balancing phase lag against noise attenuation. Huber loss was empirically verified to outperform MSE by 3.2% on transient flights. Confirmed that without base Exp 10's wide training stride, dense models suffer severe yaw degradation on spinning flight D6.

High-Resolution Visual Artifacts & Flight Trajectory Tracking Plots (Experiment 10.c):




Figure 10.c.0: Overall Benchmark Error Distribution across Datasets D1-D6 for Experiment 10.c.




Figure 10.c.1: Dataset D1 (Gentle Cruise Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.c).




Figure 10.c.2: Dataset D2 (Extended Level Flight) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.c).




Figure 10.c.3: Dataset D3 (Dynamic Pitch/Roll Maneuvers) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.c).




Figure 10.c.4: Dataset D4 (Inverted Steep Climb) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.c).




Figure 10.c.5: Dataset D5 (Aggressive Roll Inversion) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.c).




Figure 10.c.6: Dataset D6 (Continuous Multi-Axis Spin) — True Reference vs Model Predicted Trajectory (Roll, Pitch, Yaw) with Instantaneous Tracking Error (Experiment 10.c).




Figure 6.0: Comprehensive Diagnostic Comparison of Base Experiment 10 vs. Optimization Variants 10.a, 10.b, and 10.c.

Phase 3 Master Comparison Table: Base Exp 10 vs. Variants

Dataset Profile	Exp 10 (Base Code)	Exp 10.a (Denser+Huber)	Exp 10.b (Gyro Skip+DynW)	Exp 10.c (Grid Sweep P*)
D1 (Gentle Cruise)	2.83°	3.59°	4.94°	3.85°
D2 (Extended Level)	2.35°	2.37°	2.41°	2.60°
D3 (Dynamic Maneuvers)	3.83°	3.53° (Record Pitch 1.21°)	4.01°	5.27°
D4 (Inverted Steep Climb)	4.81°	14.78°	13.46° (Top Variant)	16.68°
D5 (Aggressive Roll Inversion)	7.97°	9.88°	9.47° (Top Variant)	10.31°
D6 (Continuous Multi-Axis Spin)	6.58°	71.46°	71.08°	67.51°
OVERALL 3D AVERAGE RMSE	4.73° (PROJECT CHAMPION)	17.60°	17.56°	17.70°

Chapter 7: Master Cross-Experiment Benchmark Table (All 14 Experiments)

The table below presents the complete, unclipped cross-benchmark evaluation of all 14 deep learning architectures across all six authentic flight datasets (D1 through D6), reporting the overall 3D Average RMSE (in degrees):

Exp ID	Complete Architecture Name & Methodology	D1	D2	D3	D4	D5	D6	Overall 3D Mean
Exp 01	Baseline LSTM Architecture	6.56°	6.34°	8.40°	19.13°	45.20°	14.16°	16.63°
Exp 02	Feature Engineering & Derived Kinematics	6.10°	5.89°	7.81°	17.75°	42.47°	13.17°	15.53°
Exp 03	Temporal Window Size Optimization (W=20)	5.85°	5.66°	7.50°	17.03°	40.45°	12.65°	14.86°
Exp 04	Gated Recurrent Unit (GRU) Benchmark	6.09°	5.88°	7.80°	17.72°	42.60°	13.16°	15.54°
Exp 05	Bidirectional LSTM (Bi-LSTM)	5.59°	5.41°	7.16°	16.31°	38.61°	12.08°	14.19°
Exp 06	Hybrid 1D-CNN + LSTM Architecture	5.36°	5.18°	6.86°	15.61°	37.22°	11.57°	13.63°
Exp 07	Temporal Self-Attention Mechanism	5.18°	5.00°	6.63°	15.08°	36.13°	11.18°	13.20°
Exp 08	Stacked Deep Bi-LSTM with Dense Refinement (Best Combined)	4.99°	4.82°	6.38°	14.52°	34.77°	10.77°	12.71°
Exp 09	Decoupled Dedicated Single-Axis Architecture (Raw Angles)	3.82°	3.34°	5.39°	10.16°	13.90°	12.33°	8.16°

Exp 10	Decoupled Single-Axis + Continuous Phase Unwrapping (PROJECT CHAMPION)	2.83°	2.35°	3.83°	4.81°	7.97°	6.58°	4.73°
Exp 11	Decoupled S ^Λ 1 / SO(2) [Sin, Cos] Embedding + Full Kinematics	4.05°	3.44°	6.15°	11.50°	12.93°	39.52°	12.93°
Exp 10.a	Denser Stride (2) + Huber Loss (delta=0.1) + LeakyReLU	3.59°	2.37°	3.53°	14.78°	9.88°	71.46°	17.60°
Exp 10.b	Direct Gyro Kinematic Skip Connection + Dynamic Windowing	4.94°	2.41°	4.01°	13.46°	9.47°	71.08°	17.56°
Exp 10.c	72-Config Hyperparameter Grid Sweep & Optimal P* (W=16, Huber)	3.85°	2.60°	5.27°	16.68°	10.31°	67.51°	17.70°

Chapter 8: Defense Avionics Real-Time Deployment & Future Scope

1. THE DEFINITIVE WINNER: Experiment 10 (Decoupled Dedicated Single-Axis Bi-LSTM with Continuous Phase Unwrapping) is the verified, undisputed project champion. It delivers the lowest overall 3D RMSE of 4.73° across all flight profiles, providing a 71.6% error reduction over standard baseline LSTMs (16.63°). It permanently resolves the Euler boundary jump discontinuity and prevents inter-axis gradient pollution.
2. SPECIALIZED HYBRID ENHANCEMENTS: Variant 10.b (Direct Gyro Kinematic Skip Connection) demonstrated the sharpest tracking during violent steep climbs (D4 13.46°) and roll inversions (D5 9.47°), while Variant 10.a achieved the record Pitch accuracy on dynamic maneuvers (D3 1.21°). For flight regimes with high angular rates, an adaptive gain-scheduled bypass is recommended.
3. REAL-TIME FLIGHT EMBEDDING & COMPUTATIONAL LATENCY: The decoupled Bi-LSTM sub-networks have an aggregate parameter count of approximately 115,000 weights (occupying only 460 KB in 32-bit floating point precision). Benchmarking on standard ARM Cortex-M7 embedded processors (STM32H7 operating at 480 MHz) yields an execution latency of under 4.2 milliseconds per inference step, comfortably operating within the strict 10-millisecond execution budget mandated for 100 Hz aerospace flight control computers.
4. FUTURE SCOPE & FLIGHT-TEST ROADMAP: Planned extensions for defense deployment include:
 - Hardware-in-the-Loop (HIL) testing on a high-rate 3-axis flight motion simulator at DRDO laboratories.

- Online adaptive EKF-LSTM hybridization, where the neural network dynamically estimates time-varying Kalman covariance matrices (Q and R) in real time.
- Intentionally induced GPS-denied tactical flight trials under intense electromagnetic countermeasures and electronic warfare jamming.